

Week 2 - Data Science II

Phileas Dazeley-Gaist

11/01/2022

```
# # -----  
#  
# Simple Linear Regression  
#  
# # -----
```

```
# First, let's check out this Boston data  
head(Boston)
```

```
##      crim zn indus chas   nox   rm  age   dis rad tax ptratio lstat medv  
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296    15.3  4.98 24.0  
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242    17.8  9.14 21.6  
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242    17.8  4.03 34.7  
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222    18.7  2.94 33.4  
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222    18.7  5.33 36.2  
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222    18.7  5.21 28.7
```

```
# Cool...but what do these variables mean? Let's use our ?? command to check it out and see if  
??Boston
```

```
# So medv is a column that stands for the median value of owner-occupied homes, as measured in  
# That could be a good initial variable to check out as a dependent variable.
```

```
# Let's also check out that variable lstat - "lower status of population". This variable has s  
# but it basically means "lower socioeconomic status". Let's build a model that sees how well
```

```
# We can use the lm ("linear model") command to build a linear regression:
```

```
# lm_fit <- lm(medv ~ lstat)
```

```
# why didn't this work?
```

```
?lm
```

```
#...ahhh, we forgot to tell it which dataset we were using!
```

```
lm_fit <- lm(medv ~ lstat, data = Boston)
```

```
# Let's see what is in this new object we just created
lm_fit
```

```
##
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Coefficients:
## (Intercept)      lstat
##      34.55      -0.95
```

```
# That's not super, uh, informative. Let's check it out even further with summary()
```

```
summary(lm_fit)
```

```
##
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.168  -3.990  -1.318   2.034  24.500
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 34.55384    0.56263   61.41  <2e-16 ***
## lstat       -0.95005    0.03873  -24.53  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
# We can extract stuff from this fitted model fairly easily.
```

```
# ...we can grab the names of all the internal elements if we want:
names(lm_fit)
```

```
## [1] "coefficients" "residuals"    "effects"      "rank"
## [5] "fitted.values" "assign"        "qr"           "df.residual"
## [9] "xlevels"      "call"         "terms"        "model"
```

```
# ...and thus call them!
```

```
as.data.frame(lm_fit$residuals)
```

```
##      lm_fit$residuals
## 1      -5.822595098
## 2      -4.270389786
## 3       3.974858016
## 4       1.639304221
## 5       6.709922176
## 6      -0.904083746
## 7       0.155272588
## 8      10.739604245
## 9      10.381136279
## 10      0.592003070
## 11     -0.125331595
## 12     -3.046685955
## 13      2.071434468
## 14     -6.306433217
## 15     -6.606334510
## 16     -6.606922853
## 17     -5.202516132
## 18     -3.116616860
## 19     -3.247763934
## 20     -5.637284169
## 21     -0.983803463
## 22     -1.814658317
## 23     -1.568916977
## 24     -1.166859727
## 25     -3.468036413
## 26     -4.968526049
## 27     -3.883609950
## 28     -3.336988046
## 29     -3.993209151
## 30     -2.172249621
## 31     -0.382725484
## 32     -7.665197306
## 33      4.972026713
## 34     -4.020435238
## 35     -1.729837024
## 36     -6.457363135
## 37     -3.713777753
## 38     -5.221908047
## 39     -0.229840926
## 40      0.350372329
## 41      2.227256841
## 42     -3.355602007
```

## 43	-3.734054134
## 44	-2.785473687
## 45	-4.280869551
## 46	-5.553836978
## 47	-1.110642524
## 48	-0.092913029
## 49	9.117179710
## 50	0.236958651
## 51	-2.075677071
## 52	-5.094875473
## 53	-4.537580292
## 54	-3.144924827
## 55	-1.593110444
## 56	5.415896512
## 57	-4.372056108
## 58	0.798854068
## 59	-4.736502313
## 60	-6.194385838
## 61	-3.360691877
## 62	-4.835128211
## 63	-5.960008729
## 64	-0.528372019
## 65	6.094056418
## 66	-6.617110397
## 67	-5.425335497
## 68	-4.858441114
## 69	-4.717694839
## 70	-5.302907060
## 71	-3.969509222
## 72	-3.467353264
## 73	-6.509568447
## 74	-3.990468752
## 75	-4.012506261
## 76	-4.660399657
## 77	-3.181750115
## 78	-3.996834016
## 79	-1.630231854
## 80	-5.608391760
## 81	-1.528079798
## 82	-3.794484545
## 83	-3.369509222
## 84	-4.518970233
## 85	-1.514366096
## 86	-1.750018599
## 87	0.163793810
## 88	-4.335424334
## 89	-5.728569434
## 90	-0.438559563

## 91	-3.583906073
## 92	-4.763436179
## 93	-3.901438153
## 94	-3.654034393
## 95	-3.892818223
## 96	0.163987323
## 97	-2.380281208
## 98	8.145866900
## 99	12.637835314
## 100	4.526964620
## 101	1.895624033
## 102	-0.766962336
## 103	-5.854816249
## 104	-2.485177565
## 105	-2.739732348
## 106	0.593471977
## 107	2.674080062
## 108	-0.767645485
## 109	-3.096735309
## 110	-0.380573428
## 111	-0.503199281
## 112	-2.101339445
## 113	-0.353540855
## 114	0.382502576
## 115	-6.125825133
## 116	-1.281063064
## 117	-1.915246660
## 118	-5.568332536
## 119	0.448417688
## 120	-2.323669175
## 121	1.098368334
## 122	-0.696636601
## 123	2.980544033
## 124	6.886913200
## 125	0.948026760
## 126	0.916390050
## 127	7.044504504
## 128	-2.022492488
## 129	-1.932581325
## 130	-2.829935731
## 131	-3.383219022
## 132	-3.306235802
## 133	-0.989292066
## 134	-1.874599092
## 135	-2.508486566
## 136	-0.341003840
## 137	-1.098006801
## 138	-3.592620808

## 139	-0.998788657
## 140	0.784070191
## 141	2.399351507
## 142	12.537357383
## 143	4.326482788
## 144	6.146463047
## 145	5.073104692
## 146	5.657531155
## 147	-3.135519139
## 148	8.101116537
## 149	10.151556819
## 150	1.224717759
## 151	0.341855009
## 152	-2.337185461
## 153	-7.739242712
## 154	-0.152561584
## 155	-3.189094651
## 156	-4.684099586
## 157	-6.120044310
## 158	11.106885654
## 159	-4.145023535
## 160	-4.232976155
## 161	-2.328569434
## 162	17.089744503
## 163	17.270253880
## 164	18.600322975
## 165	-0.795266402
## 166	-0.233856719
## 167	18.961341730
## 168	0.779758275
## 169	-0.208293053
## 170	-1.499282195
## 171	-3.444628705
## 172	-4.024747154
## 173	2.502384127
## 174	-2.365394721
## 175	-2.795365109
## 176	-0.090077824
## 177	-1.748841913
## 178	-3.978030444
## 179	1.920500649
## 180	7.434407864
## 181	12.428532235
## 182	10.624125514
## 183	7.925397006
## 184	3.342439450
## 185	5.127849086
## 186	7.539308123

## 187	19.673878745
## 188	3.792488804
## 189	-0.421615826
## 190	5.466925137
## 191	7.291410825
## 192	0.401890590
## 193	4.572800766
## 194	1.324907370
## 195	-1.292624710
## 196	18.267805701
## 197	2.622360484
## 198	3.926084056
## 199	6.335485842
## 200	4.678384174
## 201	2.573878745
## 202	-3.394974181
## 203	10.700812611
## 204	17.565847158
## 205	18.182301259
## 206	-1.626804404
## 207	0.268200531
## 208	5.104050449
## 209	3.773882647
## 210	7.382798699
## 211	3.553511460
## 212	7.528342624
## 213	3.075450261
## 214	2.457622059
## 215	17.220117524
## 216	-0.556873499
## 217	1.581325890
## 218	3.352137359
## 219	3.971043540
## 220	-1.578322665
## 221	1.371138346
## 222	7.534218252
## 223	2.380149203
## 224	2.766534209
## 225	14.179363445
## 226	19.844887629
## 227	6.019813598
## 228	3.088473011
## 229	15.870352587
## 230	0.518344691
## 231	0.814234092
## 232	2.133918228
## 233	9.492781024
## 234	17.498854068

## 235	2.094056418
## 236	-0.217303910
## 237	-0.390370045
## 238	1.439892564
## 239	-4.811526989
## 240	-4.251977142
## 241	-1.742279234
## 242	-2.673228893
## 243	-1.694287130
## 244	-5.923084733
## 245	-5.078223957
## 246	1.484070191
## 247	-1.551388799
## 248	-4.410839939
## 249	-1.009371032
## 250	-2.121517119
## 251	-4.548549692
## 252	-6.343163699
## 253	-1.600166661
## 254	11.609333833
## 255	-6.412016625
## 256	-4.865884357
## 257	12.400812611
## 258	20.310411812
## 259	8.847043586
## 260	2.101499662
## 261	8.357132423
## 262	15.443517429
## 263	19.860950801
## 264	7.134214350
## 265	9.641558886
## 266	-1.825825133
## 267	10.197389063
## 268	22.514526313
## 269	11.948315078
## 270	-0.885667201
## 271	-1.103199281
## 272	-3.093015638
## 273	-2.809959375
## 274	6.897483868
## 275	1.199833339
## 276	0.277306195
## 277	4.393957711
## 278	2.498364432
## 279	1.377013974
## 280	5.153898486
## 281	14.418344691
## 282	5.206885654

## 283	14.305807675
## 284	18.448315078
## 285	5.104046548
## 286	-4.734934698
## 287	-2.169702735
## 288	-4.570488494
## 289	-5.033465791
## 290	-0.718871525
## 291	-2.890176531
## 292	6.128334820
## 293	-2.188608917
## 294	-2.502417424
## 295	-2.973327600
## 296	0.002968569
## 297	-0.432976155
## 298	0.794940884
## 299	-7.332095591
## 300	-1.050606943
## 301	-3.987041302
## 302	-3.528372019
## 303	0.083087018
## 304	3.163398980
## 305	8.130001142
## 306	2.330099850
## 307	4.992978439
## 308	0.800030754
## 309	-7.440616813
## 310	-4.781848822
## 311	-6.445217048
## 312	-6.772545744
## 313	-4.019262453
## 314	-5.448450985
## 315	-1.937382877
## 316	-7.428273311
## 317	0.660563775
## 318	0.389945820
## 319	-1.611329574
## 320	-1.459712606
## 321	-3.913485532
## 322	-4.927001819
## 323	-6.838460855
## 324	-4.900261466
## 325	-3.739538834
## 326	-5.127590162
## 327	-5.711037354
## 328	-0.202709645
## 329	-5.781848822
## 330	-4.980478623

## 331	-6.117892254
## 332	-5.644727412
## 333	-7.714954439
## 334	-6.957560550
## 335	-7.441007742
## 336	-5.843945556
## 337	-5.743357213
## 338	-6.021319704
## 339	-5.868920879
## 340	-6.300360174
## 341	-7.027882383
## 342	3.361930073
## 343	-9.835913969
## 344	-3.832486519
## 345	1.025886641
## 346	-7.049821184
## 347	-5.316715567
## 348	-5.411526989
## 349	-4.363045250
## 350	-2.358050186
## 351	-5.972545744
## 352	-5.238069927
## 353	-8.552956414
## 354	-0.178618787
## 355	-8.705943582
## 356	-8.662065979
## 357	-0.032972253
## 358	-0.246685955
## 359	-0.947274298
## 360	0.083284433
## 361	-2.152956414
## 362	-1.172640550
## 363	-4.072837965
## 364	-3.845118340
## 365	-7.628079798
## 366	-0.289489481
## 367	0.646850073
## 368	1.210317006
## 369	18.543320014
## 370	18.989843210
## 371	18.258305208
## 372	24.500129462
## 373	23.882597382
## 374	12.279375151
## 375	15.319533083
## 376	-6.785177565
## 377	1.425306102
## 378	-1.074792606

## 379	1.052828311
## 380	-3.661765955
## 381	-7.803491501
## 382	-3.626800502
## 383	-0.832676131
## 384	1.079371249
## 385	3.346170826
## 386	1.917179710
## 387	2.813554845
## 388	3.238237947
## 389	4.736670333
## 390	-3.245311854
## 391	-3.198496437
## 392	6.469084997
## 393	-0.456573475
## 394	-6.341592183
## 395	-6.320533945
## 396	-5.188995943
## 397	-3.651384897
## 398	-7.128857753
## 399	-0.491831148
## 400	0.219138253
## 401	-3.521019679
## 402	-8.048838011
## 403	-3.158338505
## 404	-7.471365156
## 405	-0.041489573
## 406	-7.721706730
## 407	-0.479688963
## 408	4.870257782
## 409	7.727462060
## 410	11.738135338
## 411	-9.948841913
## 412	2.806206407
## 413	15.999355409
## 414	0.823150144
## 415	7.578984223
## 416	0.245092847
## 417	-2.552068046
## 418	1.155473905
## 419	-6.163823205
## 420	-4.549718575
## 421	-3.584099586
## 422	-5.438066025
## 423	-0.358144991
## 424	0.972808570
## 425	-6.550993969
## 426	-3.082137141

## 427	-9.447566519
## 428	-9.859124263
## 429	-3.108778787
## 430	-2.176652441
## 431	-3.294970279
## 432	-1.747369104
## 433	-7.024747154
## 434	-4.844040361
## 435	-8.441592183
## 436	0.953807583
## 437	-7.805450044
## 438	-0.725035472
## 439	6.166838135
## 440	-0.016711665
## 441	-3.048249668
## 442	1.091122506
## 443	-0.392522101
## 444	-1.245410561
## 445	-1.152166753
## 446	0.028342624
## 447	-2.752462876
## 448	-6.335029504
## 449	-3.229446096
## 450	-3.208387858
## 451	-4.584980150
## 452	-2.509465837
## 453	-2.046488540
## 454	-0.850014697
## 455	-1.878417471
## 456	-3.229446096
## 457	-3.793402664
## 458	-4.960004827
## 459	-4.234539868
## 460	-0.588115379
## 461	-2.554030491
## 462	-2.935617847
## 463	-1.762650420
## 464	-4.577833029
## 465	-0.594188423
## 466	-1.229643511
## 467	0.739505538
## 468	4.801211343
## 469	1.770553904
## 470	-0.431112418
## 471	0.822463093
## 472	-2.726705697
## 473	2.288867841
## 474	6.323734585

```
## 475      -3.519945602
## 476       1.642348546
## 477      -0.106918951
## 478       1.111888523
## 479      -2.824451031
## 480      -0.698693852
## 481      -1.350310820
## 482      -3.500458881
## 483      -2.893994910
## 484      -2.854326613
## 485      -1.280182500
## 486      -3.302318717
## 487      -1.222101560
## 488      -3.075775779
## 489      -2.195949551
## 490      -4.781157870
## 491       1.743623940
## 492      -3.786449057
## 493      -1.770682007
## 494      -1.343748141
## 495       2.857329838
## 496       5.267027747
## 497       5.230202459
## 498      -2.858144991
## 499      -1.079203229
## 500      -2.708095638
## 501      -4.139633640
## 502      -2.966863629
## 503      -5.327392747
## 504      -5.295562524
## 505      -6.397521067
## 506     -15.167451972
```

```
# We can also extract the coefficients super easily
coef(lm_fit)
```

```
## (Intercept)      lstat
##  34.5538409   -0.9500494
```

```
# and the confidence intervals!
confint(lm_fit)
```

```
##              2.5 %      97.5 %
## (Intercept) 33.448457 35.6592247
## lstat       -1.026148 -0.8739505
```

```
# most importantly, we can PREDICT stuff very, very easily with this fitted model.
# The predict() function can be used to produce confidence intervals and
# prediction intervals for the prediction of medv for a given value of lstat.
```

```
predict(lm_fit, data.frame(lstat = c(5, 10, 15))),
       interval = "confidence")
```

```
##          fit      lwr      upr
## 1 29.80359 29.00741 30.59978
## 2 25.05335 24.47413 25.63256
## 3 20.30310 19.73159 20.87461
```

```
predict(lm_fit, data.frame(lstat = c(5, 10, 15))),
       interval = "prediction")
```

```
##          fit      lwr      upr
## 1 29.80359 17.565675 42.04151
## 2 25.05335 12.827626 37.27907
## 3 20.30310  8.077742 32.52846
```


```
# # -----
#
# STOP! Your turn. Predict the value of medv for a given lstat value of 8.
# Explain the difference
#
# # -----
```

```
predict(lm_fit, data.frame(lstat = c(8)))
```

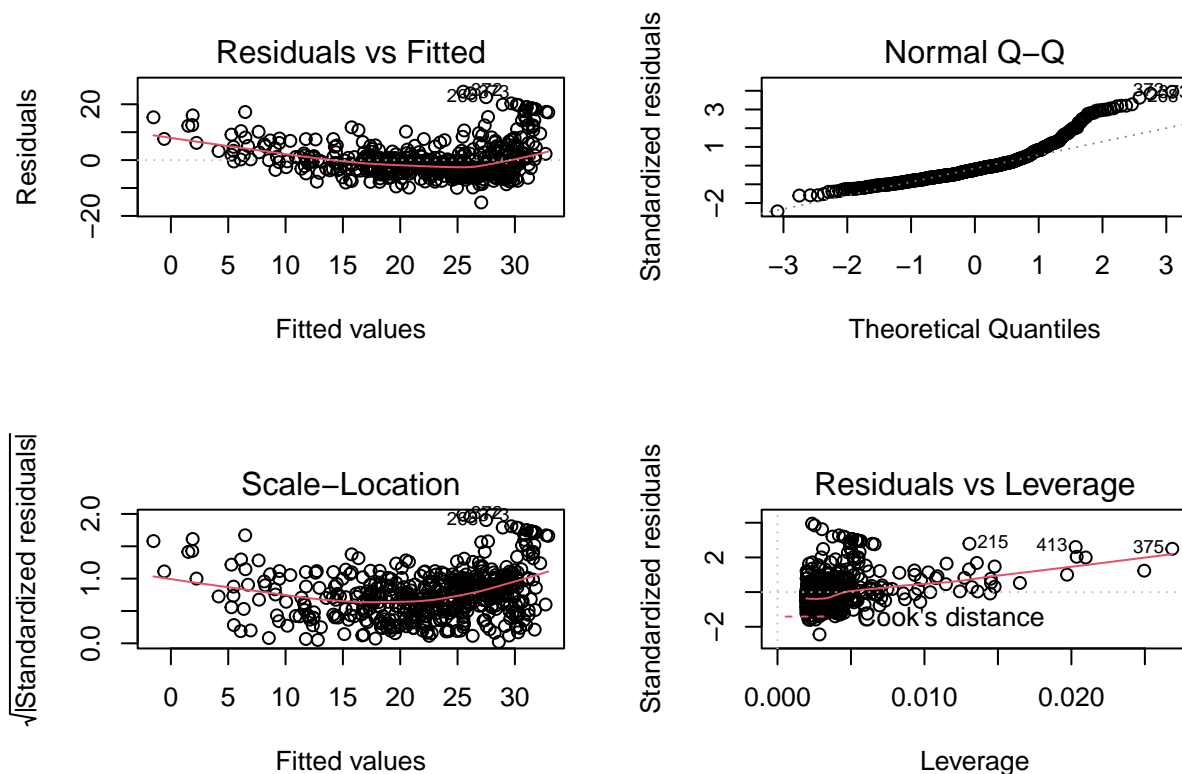
```
##          1
## 26.95345
```

```
# Let's plot these two variables and add the fitted line! We won't use ggplot2 here at the moment
# as standard plotting will work okay.
```

```
plot(Boston$lstat, Boston$medv)
abline(lm_fit, col =)
```



```
# Let's take a peek at some diagnostic plots.  
# Four diagnostic plots are automatically produced by applying the plot() function directly to  
# command will produce one plot at a time, and hitting Enter will generate  
# the next plot. However, it is often convenient to view all four plots together.  
# We can achieve this in base R graphics by using the par() and mfrow() functions, which tell  
# to split the display screen into separate panels so that multiple plots can  
# be viewed simultaneously. For example, par(mfrow = c(2, 2)) divides the  
# plotting region into a 2 x 2 grid of panels.  
par(mfrow = c(2, 2))  
plot(lm_fit)
```



What are these diagnostics?

Residuals vs. Fitted

This plot shows if residuals have non-linear patterns.

*# There could be a non-linear relationship between predictor variables and an outcome variable
show up in this plot if the model doesn't capture the non-linear relationship. If you find e
around a horizontal line without distinct patterns, that is a good indication you don't have*

Normal Q-Q

*# This plot shows if residuals are normally distributed. Do residuals follow a straight line w
It's good if residuals are lined well on the straight dashed line.*

Scale-Location

*# It's also called Spread-Location plot. This plot shows if residuals are spread equally along
This is how you can check the assumption of equal variance (homoscedasticity). It's good if
(randomly) spread points.*

Residuals vs. Leverage

*# This plot helps us to find influential cases (i.e., subjects) if any. Not all outliers are i
(whatever outliers mean). Even though data have extreme values, they might not be influential
the results wouldn't be much different if we either include or exclude them from analysis. T
and they don't really matter; they are not influential. On the other hand, some cases could
within a reasonable range of the values. They could be extreme cases against a regression li
them from analysis. Another way to put it is that they don't get along with the trend in the*


```

# Unlike the other plots, this time patterns are not relevant. We watch out for outlying values
# right corner. Those spots are the places where cases can be influential against a regression
# Cook's distance. When cases are outside of the Cook's distance (meaning they have high Cook's
# the regression results. The regression results will be altered if we exclude those cases.

# Let's keep going, so let's turn the plotting function off!
dev.off()

```

```

## null device
##          1

```

```

# # -----
#
# Multiple Linear Regression
#
# # -----

# Let's add age!
lm_fit_multiple <- lm(medv ~ lstat + age, data = Boston)
summary(lm_fit_multiple)

```

```

##
## Call:
## lm(formula = medv ~ lstat + age, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.981  -3.978  -1.283   1.968  23.158
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  33.22276    0.73085  45.458 < 2e-16 ***
## lstat       -1.03207    0.04819 -21.416 < 2e-16 ***
## age          0.03454    0.01223   2.826  0.00491 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.173 on 503 degrees of freedom
## Multiple R-squared:  0.5513, Adjusted R-squared:  0.5495
## F-statistic:  309 on 2 and 503 DF,  p-value: < 2.2e-16

```

```

# Heck, let's add EVERYTHING!
lm_fit_multiple <- lm(medv ~ ., data = Boston)
summary(lm_fit_multiple)

```

```

##

```

```
## Call:
## lm(formula = medv ~ ., data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.1304  -2.7673  -0.5814   1.9414  26.2526
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  41.617270   4.936039   8.431 3.79e-16 ***
## crim        -0.121389   0.033000  -3.678 0.000261 ***
## zn           0.046963   0.013879   3.384 0.000772 ***
## indus        0.013468   0.062145   0.217 0.828520
## chas         2.839993   0.870007   3.264 0.001173 **
## nox        -18.758022   3.851355  -4.870 1.50e-06 ***
## rm           3.658119   0.420246   8.705 < 2e-16 ***
## age          0.003611   0.013329   0.271 0.786595
## dis         -1.490754   0.201623  -7.394 6.17e-13 ***
## rad          0.289405   0.066908   4.325 1.84e-05 ***
## tax         -0.012682   0.003801  -3.337 0.000912 ***
## ptratio     -0.937533   0.132206  -7.091 4.63e-12 ***
## lstat       -0.552019   0.050659 -10.897 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.798 on 493 degrees of freedom
## Multiple R-squared:  0.7343, Adjusted R-squared:  0.7278
## F-statistic: 113.5 on 12 and 493 DF, p-value: < 2.2e-16

# is this better than our simple model?
summary(lm_fit)
```

```
##
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.168  -3.990  -1.318   2.034  24.500
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  34.55384   0.56263   61.41 <2e-16 ***
## lstat       -0.95005   0.03873  -24.53 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
# We can test if it is SIGNIFICANTLY better with an F-test!
# Use the anova function, and use the form anova(simpler_model, more complicated_model).
# Order matters!
anova(lm_fit, lm_fit_multiple)
```

```
## Analysis of Variance Table
##
## Model 1: medv ~ lstat
## Model 2: medv ~ crim + zn + indus + chas + nox + rm + age + dis + rad +
##      tax + ptratio + lstat
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      504 19472
## 2      493 11349 11      8123 32.077 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# We can check if our variables are highly collinear by using the vif()
# command from the car package.

vif(lm_fit_multiple)
```

```
##      crim      zn      indus      chas      nox      rm      age      dis
## 1.767486 2.298459 3.987181 1.071168 4.369093 1.912532 3.088232 3.954037
##      rad      tax ptratio      lstat
## 7.445301 9.002158 1.797060 2.870777
```

```
# A rule of thumb commonly used in practice is if a VIF is > 10, you have high multicollinearity.
# In our case, with values around 1-3, we are in good shape, and can proceed with our regression.
```

```
# # -----
#
# STOP! Your turn. Generate a model that predicts medv on
# the variables crim, age, and tax. Report the Adjusted R-Squared for
# this model.
#
# # -----

lm_fit_multiple_2 <- lm(medv ~ crim + age + tax, data = Boston)
summary(lm_fit)
```

```
##
```

```
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.168  -3.990  -1.318   2.034  24.500
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 34.55384    0.56263   61.41  <2e-16 ***
## lstat       -0.95005    0.03873  -24.53  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
# # -----
#
# Interaction terms and non-linear transformations!
#
# # -----

# Perhaps we sometimes want to add in the interaction effects between variables.
# You can specify an interaction term by using either the * or : character between variables.
# : generates JUST the interaction, whereas * generates the whole list (variables + interaction)

summary(lm(medv ~ lstat * age, data = Boston))
```

```
##
## Call:
## lm(formula = medv ~ lstat * age, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.806  -4.045  -1.333   2.085  27.552
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 36.0885359  1.4698355  24.553  < 2e-16 ***
## lstat       -1.3921168  0.1674555  -8.313 8.78e-16 ***
## age         -0.0007209  0.0198792  -0.036  0.9711
## lstat:age     0.0041560  0.0018518   2.244  0.0252 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 6.149 on 502 degrees of freedom
## Multiple R-squared:  0.5557, Adjusted R-squared:  0.5531
## F-statistic: 209.3 on 3 and 502 DF,  p-value: < 2.2e-16
```

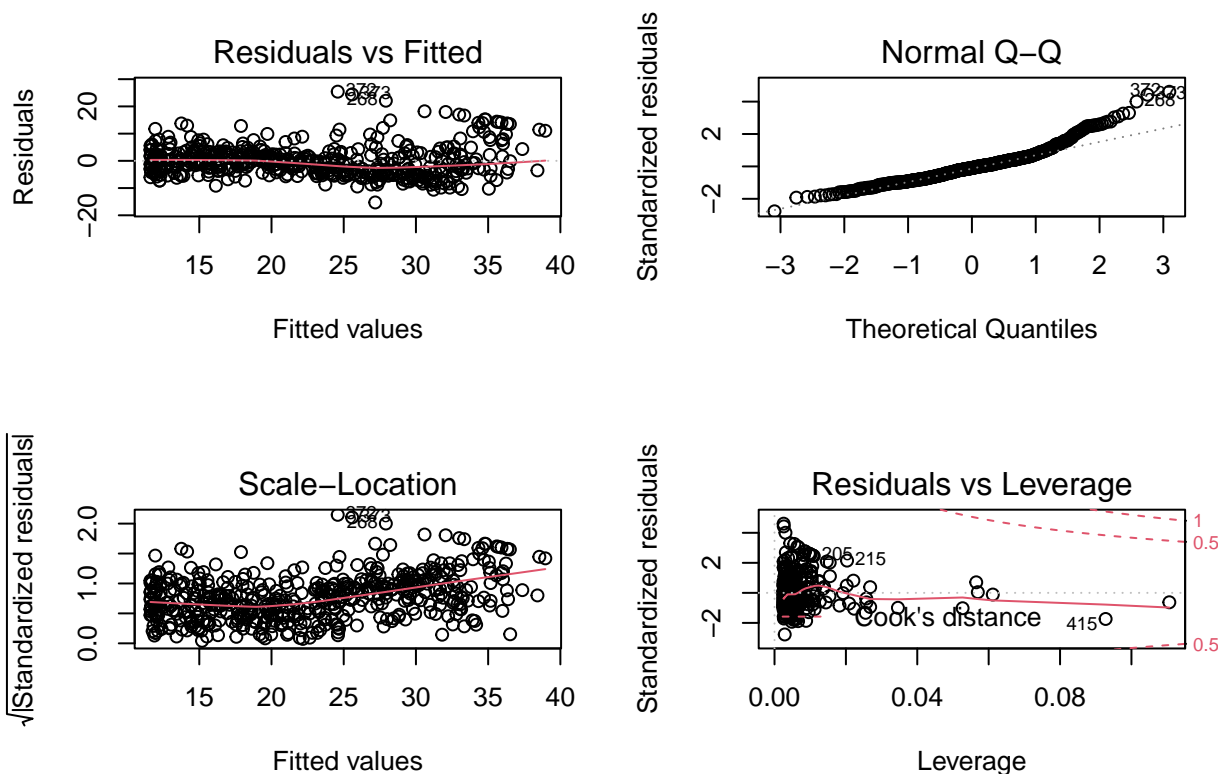
```
# The lm() function can also accommodate non-linear transformations of the
# predictors even though it's building a linear model. For instance, given a predictor X,
# we can (for example) create a predictor X^2 using I(X^2).
# The function I() is needed since the ^ has a special meaning I() in a formula object;
# wrapping as we do allows the standard usage in R,
# which is to raise X to the power 2.
```

```
# We now perform a regression of medv onto lstat and lstat2.
lm_fit_fancy <- lm(medv ~ lstat + I(lstat^2), data=Boston)
summary(lm_fit_fancy)
```

```
##
## Call:
## lm(formula = medv ~ lstat + I(lstat^2), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.2834  -3.8313  -0.5295   2.3095  25.4148
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  42.862007   0.872084   49.15  <2e-16 ***
## lstat        -2.332821   0.123803  -18.84  <2e-16 ***
## I(lstat^2)    0.043547   0.003745   11.63  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.524 on 503 degrees of freedom
## Multiple R-squared:  0.6407, Adjusted R-squared:  0.6393
## F-statistic: 448.5 on 2 and 503 DF,  p-value: < 2.2e-16
```

```
# # -----
#
# STOP! Your turn. Plot the diagnostics for the model with the quadratic term.
# Discuss what's different about them.
#
# # -----
#

par(mfrow = c(2, 2))
plot(lm_fit_fancy)
```



We can add an arbitrary number of polynomial variables. Let's add a 5th order polynomial!

```
lm_fit5 <- lm(medv ~ poly(lstat, 5), data=Boston)
summary(lm_fit5)
```

```
##
## Call:
## lm(formula = medv ~ poly(lstat, 5), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.5433  -3.1039  -0.7052   2.0844  27.1153
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    22.5328     0.2318  97.197 < 2e-16 ***
## poly(lstat, 5)1 -152.4595     5.2148 -29.236 < 2e-16 ***
## poly(lstat, 5)2   64.2272     5.2148  12.316 < 2e-16 ***
## poly(lstat, 5)3  -27.0511     5.2148  -5.187 3.10e-07 ***
## poly(lstat, 5)4   25.4517     5.2148   4.881 1.42e-06 ***
## poly(lstat, 5)5  -19.2524     5.2148  -3.692 0.000247 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.215 on 500 degrees of freedom
## Multiple R-squared:  0.6817, Adjusted R-squared:  0.6785
```

```
## F-statistic: 214.2 on 5 and 500 DF,  p-value: < 2.2e-16
```

```
# We can also take the log of a given variable. If you know me, you know I love logs.
summary(lm(medv ~ log(rm), data = Boston))
```

```
##
## Call:
## lm(formula = medv ~ log(rm), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.487  -2.875  -0.104   2.837  39.816
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -76.488      5.028  -15.21  <2e-16 ***
## log(rm)       54.055      2.739   19.73  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.915 on 504 degrees of freedom
## Multiple R-squared:  0.4358, Adjusted R-squared:  0.4347
## F-statistic: 389.3 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
# # -----
#
# Qualitative Predictors
#
# # -----

#We're going to switch to a different data set now - one about carseats!
head(Carseats)
```

```
##   Sales CompPrice Income Advertising Population Price ShelveLoc Age Education
## 1  9.50      138     73          11         276   120       Bad   42         17
## 2 11.22      111     48          16         260    83       Good  65         10
## 3 10.06      113     35          10         269    80      Medium  59         12
## 4  7.40      117    100           4         466    97      Medium  55         14
## 5  4.15      141     64           3         340   128       Bad   38         13
## 6 10.81      124    113          13         501    72       Bad   78         16
##   Urban  US
## 1   Yes  Yes
## 2   Yes  Yes
## 3   Yes  Yes
## 4   Yes  Yes
## 5   Yes   No
## 6   No   Yes
```

```
# Let's built a model that fits to every variable, the interaction effects between Income  
# and Advertising, and the interaction effects between Price and Age.
```

```
careseats_fit <- lm(Sales ~ . + Income:Advertising + Price:Age,  
                    data = Carseats)  
summary(careseats_fit)
```

```
##  
## Call:  
## lm(formula = Sales ~ . + Income:Advertising + Price:Age, data = Carseats)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.9208 -0.7503  0.0177  0.6754  3.3413   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)    6.5755654   1.0087470    6.519 2.22e-10 ***  
## CompPrice      0.0929371   0.0041183   22.567 < 2e-16 ***  
## Income         0.0108940   0.0026044    4.183 3.57e-05 ***  
## Advertising    0.0702462   0.0226091    3.107 0.002030 **  
## Population     0.0001592   0.0003679    0.433 0.665330   
## Price         -0.1008064   0.0074399  -13.549 < 2e-16 ***  
## ShelveLocGood  4.8486762   0.1528378   31.724 < 2e-16 ***  
## ShelveLocMedium 1.9532620   0.1257682   15.531 < 2e-16 ***  
## Age           -0.0579466   0.0159506   -3.633 0.000318 ***  
## Education     -0.0208525   0.0196131   -1.063 0.288361   
## UrbanYes       0.1401597   0.1124019    1.247 0.213171   
## USYes         -0.1575571   0.1489234   -1.058 0.290729   
## Income:Advertising 0.0007510  0.0002784    2.698 0.007290 **  
## Price:Age      0.0001068   0.0001333    0.801 0.423812   
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 1.011 on 386 degrees of freedom  
## Multiple R-squared:  0.8761, Adjusted R-squared:  0.8719   
## F-statistic: 210 on 13 and 386 DF, p-value: < 2.2e-16
```

```
# Notice how each "level" of the variable for ShelveLoc is present except for one?  
# R has created a ShelveLocGood dummy variable that takes on a value of  
# 1 if the shelving location is good, and 0 otherwise. It has also created a  
# ShelveLocMedium dummy variable that equals 1 if the shelving location is  
# medium, and 0 otherwise. A bad shelving location corresponds to a zero  
# for each of the two dummy variables. The fact that the coefficient for  
# ShelveLocGood in the regression output is positive indicates that a good  
# shelving location is associated with high sales (relative to a bad location).  
# And ShelveLocMedium has a smaller positive coefficient, indicating that a  
# medium shelving location is associated with higher sales than a bad shelving location
```



```
# but lower sales than a good shelving location.
```

```
# You can see that by checking this out:  
contrasts(Carseats$ShelveLoc)
```

```
##           Good Medium  
## Bad           0      0  
## Good          1      0  
## Medium        0      1
```

```
# # -----  
#  
# Let's take a look at manually predicting things and calculating  
# RMSE  
#  
# # -----
```

```
# Let's make up some pretend observations:
```

```
observed <- c(0.22, 0.83, -0.12, 0.89, -0.23, -1.30, -0.15, -1.4,  
             + 0.62, 0.99, -0.18, 0.32, 0.34, -0.30, 0.04, -0.87,  
             + 0.55, -1.30, -1.15, 0.20)
```

```
# Now, some pretend predictions:
```

```
predicted <- c(0.24, 0.78, -0.66, 0.53, 0.70, -0.75, -0.41, -0.43,  
             + 0.49, 0.79, -1.19, 0.06, 0.75, -0.07, 0.43, -0.42,  
             + -0.25, -0.64, -1.26, -0.07)
```

```
# Calculate the residuals:
```

```
residuals <- observed - predicted
```


```
# An important step in evaluating the quality of the model is to visualize  
# the results. First, a plot of the observed values against the predicted values  
# helps one to understand how well the model fits. Also, a plot of the residuals  
# versus the predicted values can help uncover systematic patterns in the model  
# predictions. Let's make some plots!
```

```
# this puts things on the same axis  
axis_range <- extendrange(c(observed, predicted))  
par(mfrow = c(1,2))  
# show Predicted vs. Observed  
plot(observed, predicted,  
     ylim = axis_range,  
     xlim = axis_range,  
     ylab = "predicted",
```

```

    xlab = "observed")
# Add a 45 degree reference line
abline(0, 1, col = "darkgrey", lty = 2)
# show predicted values versus residuals
plot(predicted, residuals, ylab = "residual", xlab = "predicted")
abline(h = 0, col = "darkgrey", lty = 2)

```



```

# turn off the par plot
dev.off()

```

```

## null device
##      1

```

```

# Calculate R2
caret::R2(predicted, observed)

```

```

## [1] 0.5170123

```

```

# Calculate RMSE
caret::RMSE(predicted, observed)

```

```

## [1] 0.5234883

```