

Week 6 - Data Science II

Phileas Dazeley-Gaist

01/02/2022

Today: Cross validation and resampling methods.

Note: The standard training data to testing data ratio is 70% to 30%. Note: The central limit theorem states that the average of averages in a data set tends to be normally distributed.

```
# Applied Data Science II - Week 6
```

```
# # -----
```

```
# Today we are going to talk about SPLINES and GAMS!
```

```
#
```

```
#
```

```
# # -----
```

```
#
```

```
# Load your libraries!
```

```
#
```

```
# # -----
```

```
library(ISLR2)
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr    0.3.4
```

```
## v tibble  3.1.6      v dplyr    1.0.7
```

```
## v tidyr   1.1.4      v stringr  1.4.0
```

```
## v readr   2.1.1      v forcats  0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()    masks stats::lag()
```

```
library(caret)
```

```
## Loading required package: lattice
```

```

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

library(splines)
library(npreg)
library(mgcv)

## Loading required package: nlme

##
## Attaching package: 'nlme'

## The following object is masked from 'package:dplyr':
##
## collapse

## This is mgcv 1.8-38. For overview type 'help("mgcv-package")'.

library(PerformanceAnalytics)

## Loading required package: xts

## Loading required package: zoo

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
## as.Date, as.Date.numeric

##
## Attaching package: 'xts'

## The following objects are masked from 'package:dplyr':
##
## first, last

##
## Attaching package: 'PerformanceAnalytics'

```

```
## The following object is masked from 'package:graphics':
##
##     legend
```

```
# # -----
#
# Splines!
#
# # -----
```


```
# Let's start by checking out the Wage dataset. It's not very fancy, but it's useful for teach
# This dataset is just wages and other data for a group of 3000 male workers in the Mid-Atlant
```

```
attach(Wage)
head(Wage)
```

```
##      year age      maritl      race      education      region
## 231655 2006  18 1. Never Married 1. White      1. < HS Grad 2. Middle Atlantic
## 86582  2004  24 1. Never Married 1. White      4. College Grad 2. Middle Atlantic
## 161300 2003  45      2. Married 1. White      3. Some College 2. Middle Atlantic
## 155159 2003  43      2. Married 3. Asian      4. College Grad 2. Middle Atlantic
## 11443  2005  50      4. Divorced 1. White      2. HS Grad 2. Middle Atlantic
## 376662 2008  54      2. Married 1. White      4. College Grad 2. Middle Atlantic
##      jobclass      health health_ins logwage      wage
## 231655 1. Industrial      1. <=Good      2. No 4.318063 75.04315
## 86582  2. Information 2. >=Very Good      2. No 4.255273 70.47602
## 161300 1. Industrial      1. <=Good      1. Yes 4.875061 130.98218
## 155159 2. Information 2. >=Very Good      1. Yes 5.041393 154.68529
## 11443  2. Information      1. <=Good      1. Yes 4.318063 75.04315
## 376662 2. Information 2. >=Very Good      1. Yes 4.845098 127.11574
```

```
# Let's take a look at a plot between wages and age...
```


```
Wage %>%
  ggplot(aes(x = age, y = wage)) +
  geom_point() +
  theme_bw()
```



*# It's got some pretty funky behavior, right? If we add a linear model to it, we can see it's not
fit really well ...*

```
Wage %>%  
  ggplot(aes(x = age, y = wage)) +  
  geom_point() +  
  theme_bw() +  
  geom_smooth(method = "lm", color = "red")
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



We can also just run the math and see that the fit probably won't be great:

```
summary(lm(wage ~ age, data = Wage))
```

```
##
## Call:
## lm(formula = wage ~ age, data = Wage)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -100.265  -25.115   -6.063   16.601  205.748
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  81.70474    2.84624   28.71  <2e-16 ***
## age          0.70728    0.06475   10.92  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 40.93 on 2998 degrees of freedom
## Multiple R-squared:  0.03827,    Adjusted R-squared:  0.03795
## F-statistic: 119.3 on 1 and 2998 DF,  p-value: < 2.2e-16
```

```
# Yep, not super great.
```

```
# So since we suspect some nonlinearity here, we can use a spline!  
# To use a spline is straightforward - you can call the bs() function (from the spline library,  
# and this lets you manually add in "knots" to your spline regression!
```

```
spline_regression <- lm(wage ~ bs(age, knots = c(25, 40, 60)), data = Wage)  
# the bs() function changes everything, it stands for B-spline analysis  
# the knots parameter is the position of the knots we would like to see
```


```
# for information on the relationship between the number of knots and the degrees of freedom, .
```

```
# and, as we can see - the fit is better!  
summary(spline_regression)
```

```
##  
## Call:  
## lm(formula = wage ~ bs(age, knots = c(25, 40, 60)), data = Wage)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -98.832 -24.537  -5.049  15.209 203.207   
##  
## Coefficients:  
##                                Estimate Std. Error t value Pr(>|t|)      
## (Intercept)                   60.494      9.460   6.394 1.86e-10 ***  
## bs(age, knots = c(25, 40, 60))1    3.980     12.538   0.317 0.750899   
## bs(age, knots = c(25, 40, 60))2   44.631      9.626   4.636 3.70e-06 ***  
## bs(age, knots = c(25, 40, 60))3   62.839     10.755   5.843 5.69e-09 ***  
## bs(age, knots = c(25, 40, 60))4   55.991     10.706   5.230 1.81e-07 ***  
## bs(age, knots = c(25, 40, 60))5   50.688     14.402   3.520 0.000439 ***  
## bs(age, knots = c(25, 40, 60))6   16.606     19.126   0.868 0.385338   
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 39.92 on 2993 degrees of freedom  
## Multiple R-squared:  0.08642,    Adjusted R-squared:  0.08459   
## F-statistic: 47.19 on 6 and 2993 DF,  p-value: < 2.2e-16
```

```
# we can also visualize it!
```

```
Wage %>%  
  ggplot(aes(x = age, y = wage)) +  
  geom_point() +  
  theme_bw() +  
  geom_smooth(method = "lm", color = "red",  
              formula = y ~ bs(x, knots = c(25, 40, 60)))
```



```
# note the somewhat funky way you have to pass these arguments into ggplot2!
```

```
# now, in the above, we SPECIFIED where to put the knots. We can achieve a similar result  
# by letting R choose for us!
```

```
# Let's show that these produce a vector with the same dimensions...
```

```
dim(bs(age, knots = c(25, 40, 60)))
```

```
## [1] 3000    6
```

```
dim(bs(age, df = 6))
```

```
## [1] 3000    6
```

```
# why 6?
```

```
# This df command produces a spline with six basis functions.
```

```
# This is because the bs() function naturally produces a cubic spline which, when it has  
# three knots, has seven degrees of freedom; six basis functions + one intercept.
```

```
# HOWEVER MANY KNOTS YOU WOULD LIKE IS: 2*knots = number of degrees of freedom.
```

```
attr(bs(age, df = 6), "knots")
```

```
## 25% 50% 75%
## 33.75 42.00 51.00
```


```
# If we just let R choose...
```

```
summary(lm(wage ~ bs(age, df = 6), data = Wage))
```

```
##
## Call:
## lm(formula = wage ~ bs(age, df = 6), data = Wage)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -99.681 -24.403  -5.202  15.441 201.413
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      56.314      7.258   7.759 1.17e-14 ***
## bs(age, df = 6)1    27.824     12.435   2.238  0.0253 *
## bs(age, df = 6)2    54.063      7.127   7.585 4.41e-14 ***
## bs(age, df = 6)3    65.828      8.323   7.909 3.62e-15 ***
## bs(age, df = 6)4    55.813      8.724   6.398 1.83e-10 ***
## bs(age, df = 6)5    72.131     13.745   5.248 1.65e-07 ***
## bs(age, df = 6)6    14.751     16.209   0.910  0.3629
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.91 on 2993 degrees of freedom
## Multiple R-squared:  0.08729,    Adjusted R-squared:  0.08546
## F-statistic: 47.71 on 6 and 2993 DF,  p-value: < 2.2e-16
```

```
# slightly better! does it look all that different?
```

```
Wage %>%
  ggplot(aes(x = age, y = wage)) +
  geom_point() +
  theme_bw() +
  geom_smooth(method = "lm", color = "red",
              formula = y ~ bs(x, knots = c(25, 40, 60))) +
  geom_smooth(method = "lm", color = "blue",
              formula = y ~ bs(x, df = 6))
```

```
# nah.
```

```
# Want to fit a spline of ANY degree (and not just a cubic one?) Use the NS function! This uses
# "natural" splines which are even more flexible:
```

```
summary(lm(wage ~ ns(age, df = 12), data = Wage))
```

```
##
```

```
## Call:
```

```
## lm(formula = wage ~ ns(age, df = 12), data = Wage)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -99.668 -24.334  -5.014  15.246 201.186
```

```
##
```


```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      57.384      6.921   8.292 < 2e-16 ***
## ns(age, df = 12)1  49.813      7.188   6.930 5.12e-12 ***
## ns(age, df = 12)2  45.386      9.318   4.871 1.17e-06 ***
## ns(age, df = 12)3  66.918      8.595   7.785 9.52e-15 ***
## ns(age, df = 12)4  59.476      8.747   6.799 1.26e-11 ***
## ns(age, df = 12)5  60.372      8.646   6.983 3.55e-12 ***
```

```
## ns(age, df = 12)6      67.180      9.009      7.457 1.15e-13 ***
## ns(age, df = 12)7      58.753      8.696      6.756 1.69e-11 ***
## ns(age, df = 12)8      57.421      8.488      6.765 1.60e-11 ***
## ns(age, df = 12)9      61.185      7.942      7.704 1.78e-14 ***
## ns(age, df = 12)10     55.098      7.848      7.020 2.73e-12 ***
## ns(age, df = 12)11     67.378     17.242      3.908 9.52e-05 ***
## ns(age, df = 12)12      7.282     12.226      0.596      0.551
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.89 on 2987 degrees of freedom
## Multiple R-squared:  0.08973,    Adjusted R-squared:  0.08607
## F-statistic: 24.54 on 12 and 2987 DF,  p-value: < 2.2e-16
```

Wage %>%

```
ggplot(aes(x = age, y = wage)) +
  geom_point() +
  theme_bw() +
  geom_smooth(method = "lm", color = "red",
             formula = y ~ bs(x, knots = c(25, 40, 60))) +
  geom_smooth(method = lm, color = "green",
             formula = y ~ ns(x, df = 12))
```



*# But why have any silly constraints at all! Let's just use a SMOOTHING spline.
 # Note: your textbook uses the smooth.spline function. We're going to use a different one
 # from the npreg library as it gives you way more flexibility.*

```
smoooooooooth <- npreg::ss(age, wage, nknots = 16)
smoooooooooth
```

```
##
## Call:
## npreg::ss(x = age, y = wage, nknots = 16)
##
## Smoothing Parameter spar = 0.3135218 lambda = 1.097456e-05
## Equivalent Degrees of Freedom (Df) 6.441072
## Penalized Criterion (RSS) 4762520
## Generalized Cross-Validation (GCV) 1594.345
```

```
summary(smoooooooooth)
```

```
##
## Call:
## npreg::ss(x = age, y = wage, nknots = 16)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -99.540 -24.432  -5.069  15.219 202.353
##
## Approx. Signif. of Parametric Effects:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   106.21      1.446   73.424  0.00000 ***
## x              27.01      12.026    2.246  0.02478  *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approx. Signif. of Nonparametric Effects:
##              Df Sum Sq Mean Sq F value Pr(>F)
## s(x)           4.441  250650    56439   35.48      0 ***
## Residuals 2993.559 4762520    1591
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.89 on 2994 degrees of freedom
## Multiple R-squared:  0.08804, Adjusted R-squared:  0.08635
## F-statistic: 52.05 on 5.441 and 2994 DF, p-value: <2e-16
```

```
smoooooooooth$fit$knot
```

```
## [1] 18 22 26 30 34 38 42 46 50 54 58 62 66 70 74 80
```

```
# and, lastly, let's let the ss function actually choose the number of knots for us  
# through cross-validation!
```

```
final_smooth <- npreg::ss(age, wage)  
final_smooth
```

```
##  
## Call:  
## npreg::ss(x = age, y = wage)  
##  
## Smoothing Parameter spar = 0.3169286 lambda = 1.161449e-05  
## Equivalent Degrees of Freedom (Df) 6.468966  
## Penalized Criterion (RSS) 4762594  
## Generalized Cross-Validation (GCV) 1594.4
```

```
summary(final_smooth)
```

```
##  
## Call:  
## npreg::ss(x = age, y = wage)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -99.57 -24.43  -5.07   15.22  202.42   
##  
## Approx. Signif. of Parametric Effects:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)   106.23      1.444   73.555  0.0000 ***  
## x              27.09      11.920    2.273  0.0231  *   
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Approx. Signif. of Nonparametric Effects:  
##              Df  Sum Sq Mean Sq F value Pr(>F)      
## s(x)          4.469  250220   55991   35.19    0 ***  
## Residuals 2993.531  4762594   1591   
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 39.89 on 2994 degrees of freedom  
## Multiple R-squared:  0.08803, Adjusted R-squared:  0.08632   
## F-statistic: 51.73 on 5.469 and 2994 DF, p-value: <2e-16
```

```
final_smooth$fit$knot
```

```
## [1] 18 19 20 21 22 23 25 26 27 28 29 30 32 33 34 35 36 38 39 40 41 42 43 45 46  
## [26] 47 48 49 50 52 53 54 55 56 58 59 60 61 62 63 65 66 67 68 69 70 72 73 74 75  
## [51] 76 80
```

```
length(final_smooth$fit$knot)
```


```
## [1] 52
```

```
# and let's see what this looks like...
```

```
pred <- predict(final_smooth, Wage$wage)
```

```
Wage %>%
```

```
  ggplot(aes(x = age, y = wage)) +  
  geom_point() +  
  theme_bw() +  
  geom_smooth(aes(x=age, y=wage), method = "lm", color = "red",  
               formula = y ~ bs(x, knots = c(25, 40, 60))) +  
  stat_smooth(method = "gam", formula = y ~ bs(x, k = 52), color = "purple")
```



```
# # -----
#
# Stop! Back to the lecture!
#
# # -----
```

```
# # -----
#
# Let's TALK ABOUT THEM GAMS
#
# # -----
```

Go ahead and download the dataset called "pisa_data.csv" from the Google Drive.

```
pisa_data <- read_csv("w6 data//pisa_data.csv")
```

```
## Rows: 65 Columns: 11
## -- Column specification -----
## Delimiter: ","
## chr (1): Country
## dbl (10): Overall, Issues, Explain, Evidence, Interest, Support, Income, Hea...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
head(pisa_data)
```

```
## # A tibble: 6 x 11
##   Country Overall Issues Explain Evidence Interest Support Income Health Edu
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Albania      NA      NA      NA      NA      NA      NA  0.599  0.886  0.716
## 2 Argenti~    391    395    386    385    567    506  0.678  0.868  0.786
## 3 Austral~    527    535    520    531    465    487  0.826  0.965  0.978
## 4 Austria     511    505    516    505    507    515  0.835  0.944  0.824
## 5 Azerbai~    382    353    412    344    612    542  0.566  0.78   NA
## 6 Belgium     510    515    503    516    503    492  0.831  0.935  0.868
## # ... with 1 more variable: HDI <dbl>
```

*# This is a nifty little dataset we've put together for you based on education.
 # The data set has been constructed using average Science scores by country from
 # the Programme for International Student Assessment (PISA) 2006, along with GNI per capita,
 # Educational Index, Health Index, and Human Development Index from UN data.*


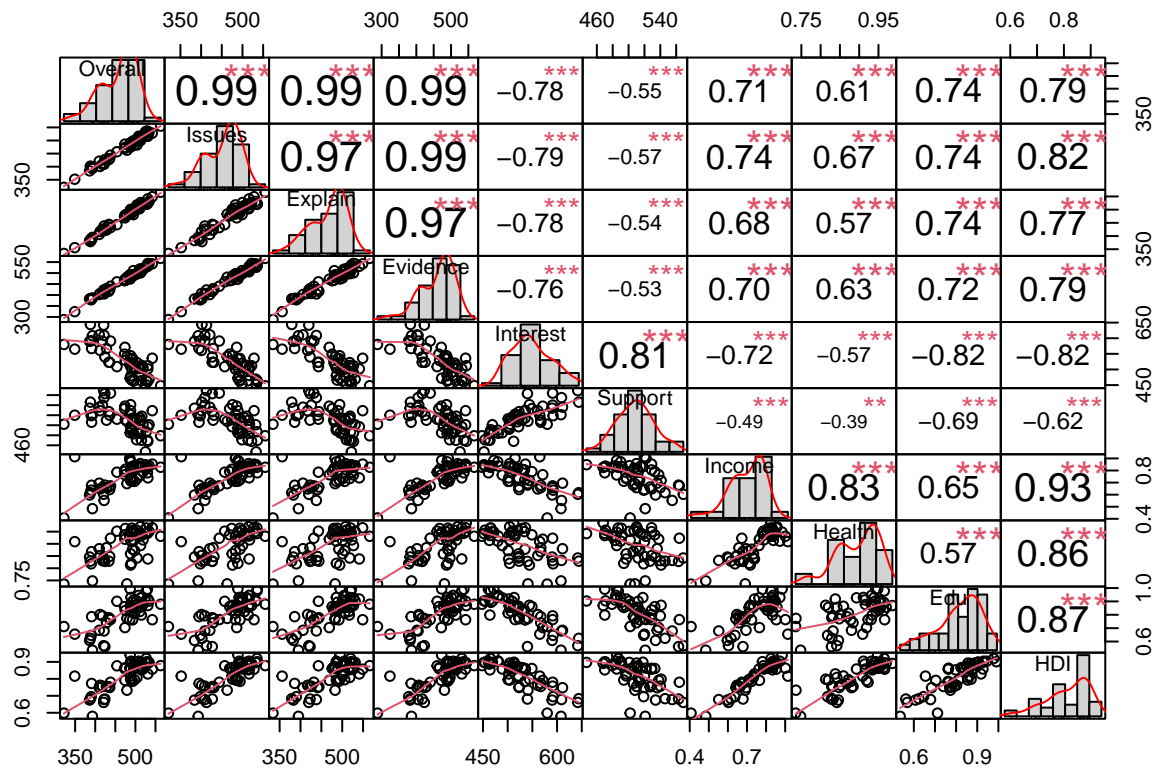
Drop nulls!

```
pisa_data <- pisa_data %>%
```

```
na.omit()
```

```
# Let's peek at the data with the chart.Correlation function
```

```
chart.Correlation(pisa_data[,2:11], histogram = TRUE, method = "pearson")
```



Simpson's paradox, also called Yule-Simpson effect, in statistics, an effect that **occurs when the marginal association between two categorical variables is qualitatively different from the partial association between the same two variables after controlling for one or more other variables.**

```
# okay, let's start simple and fit a linear model *first*. we're going to use the
# mgcv library to use the gam() function and not pass anything fancy in. Note that there are
# LOTS of libraries that use GAMs, so it's probably good to specify WHICH library you want to
# with ::
# Let's try predicting the overall score based solely on income.
```

```
pisa_lm_simple <- mgcv::gam(Overall ~ Income, data = pisa_data)
summary(pisa_lm_simple)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## Overall ~ Income
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  206.36      37.89    5.447 1.56e-06 ***
## Income       354.18      50.17    7.060 4.84e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## R-sq.(adj) =  0.489   Deviance explained = 49.9%
## GCV = 1548.4   Scale est. = 1488.9       n = 52
```

*# okay, not terribly bad. The overall adjusted R2 isn't too shabby. Now let's try fitting
a very straightforward GAM that takes advantage of splines.*

```
pisa_gam_simple <- gam(Overall ~ s(Income, bs="cr"), data = pisa_data)
```

*# Note: We again use the gam function as before for basic model fitting, but now we are using
the s function within the formula to denote the smoothing spline terms. Within that function
also specify the type of smooth, though a default is available. I chose bs = cr, denoting cu
regression splines (how we started above!)*


```
summary(pisa_gam_simple)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## Overall ~ s(Income, bs = "cr")
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  471.154      3.386   139.1  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```



```
## Approximate significance of smooth terms:
##           edf Ref.df      F p-value
## s(Income) 6.935  7.787 25.77 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.795   Deviance explained = 82.3%
## GCV = 703.74   Scale est. = 596.36      n = 52
```

```
plot(pisa_gam_simple)
```



```
# looks better! Now let's try multiple predictors - and this time, let's use a test and train
# approach to get a better idea of overall model performance.
```

```
# start linear...
```

```
set.seed(12345)
index <- createDataPartition(pisa_data$Overall, p = .8, list=FALSE)
training_data <- pisa_data[ index,]
test_data <- pisa_data[-index,]

pisa_lm_multivariate <- mgcv::gam(Overall ~ Income + Edu + Health, data = training_data)
summary(pisa_lm_multivariate)
```

```
##
## Family: gaussian
## Link function: identity
```

```
##
## Formula:
## Overall ~ Income + Edu + Health
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   48.97     100.46   0.488  0.62856
## Income        194.66     108.03   1.802  0.07909 .
## Edu           248.07      61.49   4.034  0.00024 ***
## Health         83.30     174.83   0.476  0.63632
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## R-sq.(adj) =  0.614   Deviance explained = 64.1%
## GCV = 1326.3   Scale est. = 1205.8       n = 44
```

```
# ...get rmse!
```

```
predictions_lm <- predict(pisa_lm_multivariate, test_data)
RMSE(predictions_lm, test_data$Overall)
```

```
## [1] 30.24088
```

```
# now let's go full GAMMMMMSSSSS
```

```
pisa_gam_multivariate <- gam(Overall ~ s(Income) + s(Edu) + s(Health), data = training_data)
summary(pisa_gam_multivariate)
```


```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## Overall ~ s(Income) + s(Edu) + s(Health)
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  470.318      3.048   154.3  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df    F  p-value
## s(Income)  6.948  7.904 8.025 1.03e-05 ***
## s(Edu)      5.334  6.305 2.219  0.0623 .
## s(Health)   1.000  1.000 0.953  0.3368
```


```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.869   Deviance explained =   91%
## GCV = 605.14   Scale est. = 408.72    n = 44
```

```
predictions_gam <- predict(pisa_gam_multivariate, test_data)
RMSE(predictions_gam, test_data$Overall)
```

```
## [1] 25.23021
```

```
# want to see how each of these effects is being modeled? use plot!
plot(pisa_gam_multivariate)
```





if you want it to be prettier, check out the visreg package!

```
# # -----
#
# Data Project Time!
#
# # -----
```

*# Go into the Drive and open the file "walmart.csv".
 # this is some data related to weekly store sales at a bunch of different walmart stores*

around the country. Your job is to use all the tools at your disposal (linear models, lasso models, ridge models, or GAMS) to try and have the BEST fitting model possible as based on minimizing RMSE.
You are trying to predict the weekly_sales based on the data available to you.
good luck!

```
walmart <- read_csv("w6 data/walmart.csv")
```

```
## Rows: 6435 Columns: 8
## -- Column specification -----
## Delimiter: ","
## chr (1): date
## dbl (7): store, weekly_sales, holiday_flag, temperature, fuel_price, cpi, un...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
head(walmart)
```

```
## # A tibble: 6 x 8
##   store date       weekly_sales holiday_flag temperature fuel_price   cpi
##   <dbl> <chr>         <dbl>         <dbl>         <dbl>     <dbl> <dbl>
## 1     1 05-02-2010     1643691.         0         42.3       2.57  211.
## 2     1 12-02-2010     1641957.         1         38.5       2.55  211.
## 3     1 19-02-2010     1611968.         0         39.9       2.51  211.
## 4     1 26-02-2010     1409728.         0         46.6       2.56  211.
## 5     1 05-03-2010     1554807.         0         46.5       2.62  211.
## 6     1 12-03-2010     1439542.         0         57.8       2.67  211.
## # ... with 1 more variable: unemployment <dbl>
```

```
walmart_cleaned <- walmart %>%
  mutate(store = as.factor(store),
         holiday_flag = as.factor(holiday_flag),
         year = as.factor(lubridate::year(lubridate::dmy(date))),
         month = as.factor(lubridate::month(date))) %>%
  select(-c(date))

set.seed(12345)
index <- createDataPartition(walmart_cleaned$weekly_sales, p = .8, list=FALSE)
training_data <- walmart_cleaned[ index,]
test_data <- walmart_cleaned[-index,]

walmart_model <- gam(weekly_sales ~ store + holiday_flag +
                     s(temperature) +
                     s(fuel_price) + s(cpi) + s(unemployment) +
```

```

                                year + month, data = training_data)
summary(walmart_model)

```

```

##
## Family: gaussian
## Link function: identity
##
## Formula:
## weekly_sales ~ store + holiday_flag + s(temperature) + s(fuel_price) +
##      s(cpi) + s(unemployment) + year + month
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1454624    257049   5.659 1.61e-08 ***
## store2       376933     18781  20.070 < 2e-16 ***
## store3      -1193481     21059 -56.674 < 2e-16 ***
## store4       668728     548014   1.220 0.222418
## store5      -1302932     21643 -60.200 < 2e-16 ***
## store6       -55595     20884  -2.662 0.007791 **
## store7      -723136     80315  -9.004 < 2e-16 ***
## store8      -733813     23966 -30.619 < 2e-16 ***
## store9      -1101456     24143 -45.622 < 2e-16 ***
## store10       518738     547790   0.947 0.343702
## store11      -236011     21201 -11.132 < 2e-16 ***
## store12      -293097     552588  -0.530 0.595852
## store13       614313     547736   1.122 0.262107
## store14       765706     103027   7.432 1.25e-13 ***
## store15      -847598     526213  -1.611 0.107296
## store16      -828160     80278 -10.316 < 2e-16 ***
## store17      -507798     547380  -0.928 0.353613
## store18      -374261     526884  -0.710 0.477533
## store19      -42418     526098  -0.081 0.935741
## store20       578201     27630  20.926 < 2e-16 ***
## store21      -804433     18745 -42.914 < 2e-16 ***
## store22      -470571     513460  -0.916 0.359463
## store23      -218273     525356  -0.415 0.677811
## store24      -118722     526448  -0.226 0.821587
## store25      -822434     27997 -29.376 < 2e-16 ***
## store26      -465109     526290  -0.884 0.376873
## store27       274482     513129   0.535 0.592730
## store28       16625     552728   0.030 0.976006
## store29      -924656     527766  -1.752 0.079831 .
## store30     -1116328     18912 -59.028 < 2e-16 ***
## store31      -174666     18942  -9.221 < 2e-16 ***
## store32      -129996     80188  -1.621 0.105048
## store33     -1105113     547800  -2.017 0.043710 *
## store34      -370518     549730  -0.674 0.500342

```

```

## store35      -580067      513736  -1.129 0.258903
## store36      -1181026      19108  -61.810 < 2e-16 ***
## store37      -1034095      19068  -54.230 < 2e-16 ***
## store38      -924065      552602  -1.672 0.094545 .
## store39      -94166      19261  -4.889 1.05e-06 ***
## store40      -649033      525420  -1.235 0.216789
## store41      -82595      79726  -1.036 0.300255
## store42      -809127      547749  -1.477 0.139688
## store43      -859916      35975  -23.903 < 2e-16 ***
## store44      -1105192      547493  -2.019 0.043577 *
## store45      -476212      103073  -4.620 3.93e-06 ***
## holiday_flag1 26246      8519    3.081 0.002075 **
## year2011      -94787      22025  -4.304 1.71e-05 ***
## year2012      -134851      35978  -3.748 0.000180 ***
## month2        118182      11837    9.984 < 2e-16 ***
## month3         66011      13993    4.717 2.45e-06 ***
## month4         67517      16244    4.156 3.29e-05 ***
## month5         66405      18460    3.597 0.000325 ***
## month6        106135      19786    5.364 8.50e-08 ***
## month7         76267      21366    3.570 0.000361 ***
## month8         87531      21520    4.067 4.83e-05 ***
## month9         19603      20214    0.970 0.332215
## month10        20247      18508    1.094 0.274034
## month11        156146      18928    8.249 < 2e-16 ***
## month12        307268      18794   16.349 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(temperature) 2.838  3.614 5.863 0.000261 ***
## s(fuel_price)   4.486  5.631 5.376 3.34e-05 ***
## s(cpi)          7.496  8.250 4.836 2.91e-06 ***
## s(unemployment) 4.726  5.921 7.004 8.81e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.939   Deviance explained =  94%
## GCV = 1.998e+10   Scale est. = 1.9676e+10   n = 5151

```

```

predictions_gam <- predict(walmart_model, test_data)
RMSE(predictions_gam, test_data$weekly_sales)

```

```
## [1] 138804.5
```